

RAG Knowledge Assistant — Test Document

Purpose

This document is provided as a sample knowledge source for testing the document upload workflow of an AI-powered RAG Knowledge Assistant.

What is RAG?

Retrieval-Augmented Generation (RAG) combines information retrieval with a language model. Relevant pieces of information are retrieved from a knowledge base and supplied to the model as context before generating an answer.

Typical RAG Pipeline

1. Upload a document.
2. Extract text from the document.
3. Split the text into smaller chunks.
4. Generate vector embeddings for the chunks.
5. Store the embeddings in a vector database.
6. Retrieve relevant chunks for a user question.
7. Generate an answer using the retrieved context.

Test Content

A knowledge assistant can help users search their private documents, understand technical information, summarize content, and answer questions using retrieved document context.